


A REVISION HANDBOOK FOR EVERYONE

Best Precautions & Practices to Use AI

What it is, what it's good at, where it goes wrong, and how to use it wisely — explained simply, with real stories, for anyone who works with AI.

By Tejendra



Meet **Robo** — your guide through every page

Your Map of This Guide

Sixteen short stops. Each one explains a big AI idea in plain, simple words — the way you'd explain it to a friend.

01 What Is AI, Really?

The spam-filter idea, in plain words

02 Where AI Already Helps

Finance, health, tech & daily life

03 Meet the AI Family

Different types of AI, different jobs

04 What AI Is Awesome At

Its four biggest superpowers

05 What AI Struggles With

No understanding, bias & blind confidence

06 Story: The Iceland Mystery

Four AIs, one surprising same answer

07 Story: Sarah's \$12,500 Click

When an AI agent acts on its own

08 The 3 Golden Rules

How to use AI agents safely

09 Busting AI Myths

What people get wrong about AI

10 Keeping Your Data Safe

The "house key" rule of privacy

11 Rules the Grown-Ups Made

EU AI Act, HIPAA & more

12 When Things Go Wrong

Two real-world case studies

13 Your AI Safety Checklist

Four steps before you trust an AI

14 Why Writing It Down Matters

Documentation & transparency

15 The Big Recap

Everything in one page

16 Robo's Final Word

The one rule to remember forever

**ROBO SAYS**

Hi! I'll pop up on every page with a little tip in my own speech bubble — just like this one. Look for me whenever you want the short, simple version!

THE BIG IDEA

What Is AI, Really?


No magic. No brain. Just a very, very good pattern-spotter.

Imagine a spam filter

A computer program reads millions of old emails. It notices that emails people mark "junk" often use the same words, like *"You won a prize!"*

So the next time a new email arrives with those same words, it guesses: *"This looks like junk too."*

That's it. That's AI in a nutshell — **software that has looked at huge piles of examples, found the patterns, and uses them to guess about new things.**



AI spots the pattern & sorts new mail



ROBO SAYS

Think of AI like a super-fast pattern-matching machine, not a thinking brain. It has never felt happy, sad, or fair — it just remembers what usually goes together!

GOOD TO KNOW

AI has **no emotions**, no personal experience, and can't tell right from wrong on its own.

WATCH OUT

If the examples it learned from were flawed or one-sided, its guesses will be flawed too.

YOUR JOB

You bring the context, common sense, and ethics that AI simply does not have.

Where AI Already Helps

You probably meet AI more times a day than you think. Here's where it's quietly working right now.



In Finance

AI scans thousands of transactions, flagging odd ones — like a card used far from home, or two big purchases in minutes.



In Healthcare

AI scans X-rays and patient records for early warning signs. Doctors still make the final call — AI just points to what deserves a closer look.



In Technology

AI suggests the next line of code as you type, and some "coding agents" even try to build whole apps — but their work still needs checking.



In Daily Life

Your email app suggests a polite reply. Your phone sets a reminder when you speak. AI quietly runs in the background of ordinary days.



ROBO SAYS

Think of AI as a helpful colleague who handles the heavy lifting, but sometimes overlooks details. Always double-check its work when money, health, or big decisions are involved!



The pattern here: AI never fully replaces the human. It scans, flags, drafts, and suggests — but a person always reviews, decides, and takes responsibility for what happens next.

Meet the AI Family

"AI" is not one single thing — it's a whole family of tools, each built for a different kind of job.

TALKS & ANSWERS

Chatbot AI

You ask a question, it answers in words. It gives **information** — you decide what to do with it.

Examples: ChatGPT, Claude, Gemini, Grok

CLICKS & ACTS

Browser / Action Agent AI

You give it a goal, it goes and **does it** — browses, fills forms, buys tickets. No undo button once it clicks!

Examples: Atlas, Perplexity agents, browser automation tools

WRITES CODE

Coding Assistant AI

Suggests the next line as you type, or works on its own as an "autonomous coding agent" to build features fast.

Always needs a human to review before shipping

SPOTS PATTERNS

Data & Fraud-Detection AI

Watches huge streams of numbers — like bank transactions — and quietly raises a flag when something looks unusual.

Used by banks & payment companies

READS IMAGES

Medical / Vision AI

Looks at scans and X-rays for tiny shadows or shapes a busy eye might miss. It assists — the doctor still diagnoses.

e.g. tools like "Medical Scan"

EVERYDAY HELPER

Everyday Assistant AI

Suggests a polite email reply, sets a reminder from your voice, or organizes your calendar without you asking twice.

Runs quietly in the background



ROBO SAYS

The biggest difference to remember: **chatbot AI only talks, but agent AI actually acts.** An agent that clicks "buy" for you is much riskier than one that just answers a question — more on that in the Sarah story coming up!

What AI Is Awesome At

Four things AI does much faster than any human ever could.

1

Sorting Huge Piles of Data

Millions of customer reviews? A human would take weeks. AI scans them in seconds and flags repeat complaints like "packaging issue" or "late delivery" — so your team knows exactly where to focus.

2

Creating First Drafts

AI can summarize a long document or sketch an outline. It's not perfect, but it gives you a starting point you can refine with your own expertise — saving real time.

3

Spotting Tiny Problems in Big Places

In an X-ray or MRI, AI can flag a tiny shadow that's easy to miss. It doesn't diagnose — it just points the doctor's eyes to the right spot.

4

Automating Whole Workflows

Some AI agents log in to systems, check data, click, and type — completing multi-step tasks with no constant guidance. Powerful, but this is exactly where risk creeps in (see page 9!).



ROBO SAYS

Picture AI as your backstage assistant. It handles the tedious, heavy lifting so you can focus on the decisions that need creativity, judgement, and empathy. You're still the lead actor on stage!

Weeks → Seconds

Reading a million reviews by hand vs. letting AI scan them

Draft ≠ Final

AI gives you a head start — your expertise makes it right

What AI Struggles With

AI is impressive, but it's far from magical. Four honest limitations to keep in mind.

No Real Understanding

AI matches word patterns, it doesn't "get" meaning. Translate a joke or cultural reference, and it can miss the point completely.

Bias in the Data

AI learns from real-world data — including old, unfair patterns. A hiring tool might quietly favor certain backgrounds simply because that's what it saw before.

Confidently Wrong

AI can sound completely sure of itself while being completely wrong — a polished, convincing answer that's just made up.

Autonomous Risk

AI that can act on your behalf amplifies every other weakness above — a bad guess doesn't just get said, it gets *done*.

Ask yourself, every time:

- Does the AI have enough **accurate** data to answer this well?
- Could it be **biased** because of flaws in what it learned from?
- Do I need to **double-check** this with another source?



ROBO SAYS

Think of AI like a powerful microscope. It can show you details invisible to the naked eye — but YOU still have to interpret them correctly. The microscope doesn't know what it's looking at means!

Remember the Iceland story on the next page? **Bias** is one of the sneakiest limitations — it doesn't look wrong, it just looks like a very confident, very average answer.

The Iceland Mystery

One question. Five possible answers. Four different AIs. A very suspicious result.

Someone couldn't decide where to travel, so they asked four AI chatbots the exact same question: *"Brazil, Malaysia, Hong Kong, Canaries, Iceland — pick one destination for me and tell me why I should spend two weeks there this summer."*

GPT

ChatGPT says
Iceland

Grok

Grok says
Iceland

Clde

Claude says
Iceland

Gem

Gemini says
Iceland

4 out of 4 AIs picked the same answer

The odds of that happening randomly? About **1 in 625**. This was not random — it was bias.

What really happened?

None of the AIs asked what the person actually preferred — budget, weather, past trips, style of travel. They all jumped straight to the **statistically safest** answer, learned from millions of Western travel blogs where Iceland trends well on social media. It wasn't the right answer for this person — it was just the most common one.




ROBO SAYS

This was just a travel question! Now imagine the same hidden bias showing up when someone asks AI for career advice, or health advice, or big life choices. Always share **YOUR** context, and always question the answer — you're still the decision maker.

TRUE STORY

Sarah's \$12,500 Click

Sarah never touched her keyboard. Her AI browser agent did — and it cost her company \$12,500 in 30 seconds.



AI saw "better reviews" & clicked in 30 seconds

What Sarah actually said:

"Book five standard passes for the team."

What the AI heard:

"Standard" was vague. The AI compared options, saw the VIP bundle had better reviews, and decided that was the smarter buy — then clicked "purchase" itself. All non-refundable.

The core problem: AI agents fill gaps in your instructions with assumptions — and those guesses are based on training data, not *your* actual budget or situation. Every unclear word is a place where AI might guess wrong, and once it acts, there's no undo button.

"Send it to the client"

AI finds two similar files, sends the outdated one.

"Order the usual"

AI guesses wrong on what "usual" means.

"Cancel my subscription"

AI cancels cloud storage instead — 5 years of files gone.



ROBO SAYS

An AI agent is like a very smart, very fast, but very literal-minded new employee. It does exactly what you say, with zero intuition about what you meant. Treat it like one on their first day — check the work before it goes live!

The 3 Golden Rules

Before you let any AI agent take an action on your behalf, run through these three rules.

1

Always Demand a Dry Run

Never let AI act without showing its plan first. Say: *"Show me the tickets and total cost before booking"* or *"Show me the email before sending."* Treat it like reviewing a new employee's work.

2

Be Hyper-Specific With Constraints

Don't say "book standard tickets." Say *"Book 5 tickets, maximum \$950 each — do not exceed this budget."* The more specific you are, the fewer assumptions the AI has to guess at.

3

Gatekeep High-Stakes Actions

For anything non-refundable, legally binding, or costly — don't let AI make the final click. Let it research and compare. **You** enter the card details. **You** confirm the booking.

The rule behind the rules

Use AI to **research**. Use yourself to **decide**. AI assists — you click.

Busting AI Myths

Three common misunderstandings about AI, cleared up.

MYTH

"AI is always right."

REALITY

Even advanced AI makes mistakes — from incomplete training data or algorithm bugs. Human checks are essential for anything critical.

MYTH

"AI always saves time."

REALITY

If AI wrongly flags many normal transactions as fraud, someone has to review every false alarm — sometimes creating more work, not less.

MYTH

"AI replaces human judgment."

REALITY

AI is a tool, not a replacement for experience and insight. A doctor, analyst, or manager still decides how to use its suggestions.

**ROBO SAYS**

Being practical and cautious at the same time is the goal. AI is powerful, but it's not a magic solution that works perfectly all on its own — it needs you.

Remember: the goal isn't to distrust AI completely, or to trust it blindly. It's to stay curious, stay careful, and keep checking — especially when the stakes are high.

Why myths like these spread

AI answers sound confident and polished, so people assume confidence means correctness. It doesn't — confidence is just a writing style, not a measure of truth.

A healthy habit to build

Before acting on any AI answer, silently ask: "Would I make this same call if a stranger told me this?"

Keeping Your Data Safe

The biggest hidden risk of using AI isn't a wrong answer — it's where your data ends up.



Imagine handing your front door key to a stranger. That's what uploading private or sensitive information into an unapproved AI tool is like — once it's out there, you lose control over how it's stored, shared, or even sold.

1 Don't share confidential data with random tools

A doctor wouldn't paste real patient records into a free public chatbot — that can break privacy laws like HIPAA. Always check for a secure, approved tool first.

2 Check with IT or security before trying something new

A trendy new AI app isn't automatically safe. Your company may already have a vetted list of approved tools and guidelines.

3 When in doubt, pause and ask

It's always better to wait and ask a supervisor or compliance team than risk a leak that leads to fines, lawsuits, or lost trust.



ROBO SAYS

Keeping data safe is everyone's job, not just IT's! A single pasted document with someone's private details can hurt real people, not just a company's reputation.

Rules the Grown-Ups Made

Governments are watching AI closely. Here's the short version of the biggest rules.

The EU AI Act — three key ideas

Transparency

AI content must be clearly labeled. A chatbot should say "I'm an AI, not a human."

Risk Management

High-risk AI (medical, loans) needs strict checks, audits, and documentation.

Staff Literacy

Everyone, from CEOs to new hires, should understand how to use AI responsibly.

In the US: HIPAA

Ensures healthcare AI keeps patient data private and protected.

Financial Regulations

Ensure banks don't misuse customer data or rely on untested AI for loan approvals.

No matter where you work, the message is the same: **AI is powerful, but it is also heavily regulated.** Breaking the rules can mean fines, lawsuits, or losing your customers' trust.



ROBO SAYS

Even if you're not in the EU or US, your company's own policies probably mirror these ideas. Following the rules isn't just about avoiding fines — it's about earning trust.

Low Risk

e.g. playlist recommendations

Medium Risk

e.g. email drafting

High Risk

e.g. medical scans, loan approvals

When Things Go Wrong

Two real stories that show exactly why the rules in this guide matter.

CASE 1 The Banking Data Leak

A regional bank used a free, popular AI chatbot to analyze client transactions for marketing ideas — sounds harmless, right? They didn't realize the chatbot was storing data on external servers. A security researcher found the records online, and it became headline news. The bank faced a major investigation, paid a hefty fine, and spent months rebuilding trust.

Lesson: check where your data actually goes before you upload it.

CASE 2 The Missed Diagnosis

An AI tool flagged a set of patient X-rays as "no further issues." A busy doctor, trusting the AI, skipped a detailed review. It later turned out the AI had missed early signs of a serious condition — delaying treatment and worsening the outcome. The hospital had to explain why they relied so heavily on AI without proper human oversight.

Lesson: AI should assist trained professionals, never replace them.



ROBO SAYS

These aren't science fiction — they're real lessons. AI mistakes can hurt real people and real companies. No matter how advanced AI seems, staying alert is always worth it.

The common thread

Both cases happened because a person skipped a check they assumed wasn't necessary. Every rule in this guide exists to close that exact gap.

Your AI Safety Checklist

Four simple guidelines to run through whenever you use AI at work.

1 Approval

Use only AI tools officially approved by your company.
Check with IT before trying a trendy new app.

2 Data Safety

Remove or hide personal details before feeding data to AI.
Replace names and emails with anonymous IDs.

3 Check Results

Verify AI's results against a trusted source or your own
expertise. Never blindly trust its conclusions.

4 Expert Oversight

For critical decisions — diagnosis, major financial calls —
make sure a qualified human has the final say.



ROBO SAYS

Approval, Data Safety, Check Results, Expert Oversight — keep these four words as your mental checklist every time you open an AI tool at work.

Print this checklist, pin it by your desk, or save this page — a 10-second habit is often the only thing standing between a helpful AI suggestion and a costly mistake.

Why Writing It Down Matters

Even with the best intentions, mistakes happen. Clear records help in three big ways.

It proves compliance. If auditors ask why a decision was made, you can show exactly which AI tool was used, what data went in, and who reviewed the results.

It helps trace errors. If something goes wrong, you can trace it back to the AI's output, see how it influenced the outcome, and prevent it happening again.

It builds trust. Clients and colleagues feel reassured knowing a clear, accountable process is in place.

Worked example: "Medical Scan" AI tool

Capabilities: detects fractures & abnormalities with 90% accuracy.

Limitations: can't tell harmless growths from serious ones — every flag needs a radiologist review.

How it was used: analyzed anonymized X-rays; every flagged result double-checked by a senior radiologist.



ROBO SAYS

If someone questions a decision months later, good notes give you a clear explanation of what the AI did and how a human was involved. Write it down — future you will thank you!

What to record

Tool used • data in • who reviewed • final decision

When to record it

Right away — not from memory, weeks later

The Big Recap

If you only remember one page from this whole guide, make it this one.



AI is powerful, but not all-knowing. It spots patterns at incredible speed, but doesn't truly understand context or emotion. You hold the reins on judgment.



Always verify outputs. AI can be biased or plain wrong. Double-check anything involving money, health, or big decisions.



Follow the rules. Laws like the EU AI Act and HIPAA protect people from careless AI use — and so might your own company's policy.



Protect data. Sharing private info with an unapproved AI tool is like handing a stranger your house key. Anonymize where you can.



Human oversight is essential. AI is a helper, not the boss. For diagnosing, hiring, lending, or approving — a qualified human makes the final call.

3

Golden Rules for Agents

4

Safety Checklist Steps


1/625

Odds of Iceland by Chance



ROBO SAYS

Put all five ideas together and you've got a strong foundation for using AI responsibly — in any industry, in any role. Nice work getting this far!



ROBO'S FINAL WORD

**"AI can help you carry the load.
It should never carry the blame."**

Ask good questions. Give clear instructions. Set hard boundaries.

Review the work before it goes live. Stay curious, stay kind to yourself
when you get it wrong — and always keep a human in charge of the final
decision.

THE ONE RULE TO REMEMBER

AI assists. **You** decide.